

Laboratory 13: TensorFlow image classification

In this laboratory exercise, you will use TensorFlow to train an efficient machine learning model to recognize the 10 native North American species of *Acer* (maples)¹ from herbarium specimen images. The images been processed to redact the specimen labels with opaque boxes—this prevents the model from learning to read. Each image was reshaped into a square (with distortion) and resized to 64×64 pixels. The images were randomly divided into 8,990 train (82.4%), 960 test (8.8%), and 960 validate (8.8%) partitions. There 1091 images for each species.

The machine learning model that you will train is a variant of reduced FireNet (Gupta et al. 2023) which is a smaller-scale SqueezeNet (Iandola et al. 2016). This tiny model (22,090 parameters) was selected for a combination of relatively rapid training/inference speed and its suitability for the North American *Acer* dataset.

Tasks

- (1) Install Netron by typing `sudo snap install netron` in the terminal. Type your password when prompted.
- (2) Create a TensorFlow environment to build and train the FireNet model.
 - (a) Make a working project directory by typing `mkdir acer && cd acer` in the terminal.
 - (b) Start a Docker build file by typing `echo 'FROM tensorflow/tensorflow:2.9.3' > Dockerfile` in the terminal.
 - (c) Add a RUN statement to the build file by typing `echo 'RUN python3 -m pip install --upgrade pip && python3 -m pip install --upgrade setuptools && python3 -m pip install tensorflow-addons' >> Dockerfile` in the terminal.
 - (d) Build the Docker image by typing `docker build -t 'tensorflow:2.9.3' .` in the terminal. This process may take several minutes to complete depending on your network speed.
 - (e) Download the segmented data archive by typing `for k in {0..4}; do wget 'https://github.com/dpl10/phytoinformatics2023/raw/main/NA-Acer-data'$k'.tar'; done` in the terminal. Answer question (1).
 - (f) Extract the archive by typing `tar -x -M -fNA-Acer-data{0..5}.tar` in the terminal. Three .tfr files should be output.
 - (g) Download the model script by typing `wget https://raw.githubusercontent.com/dpl10/phytoinformatics2023/main/firenet.py` in the terminal.
 - (h) Download the training script by typing `wget https://raw.githubusercontent.com/dpl10/phytoinformatics2023/main/trainImages.py` in the terminal.
 - (i) Download the testing script by typing `wget https://raw.githubusercontent.com/dpl10/phytoinformatics2023/main/testImages.py` in the terminal.
 - (j) Download the fine-tune training script by typing `wget https://raw.githubusercontent.com/dpl10/phytoinformatics2023/main/finetuneImages.py` in the terminal.

¹ *A. circinatum* Pursh, *A. glabrum* Torr., *A. grandidentatum* Nutt., *A. macrophyllum* Pursh, *A. negundo* L., *A. pensylvanicum* L., *A. rubrum* L., *A. saccharinum* L., *A. saccharum* Marshall, and *A. spicatum* Lam.

- (k) Download a pretrained FireNet model by typing
`wget https://github.com/dpl10/phytoinformatics2023/raw/main/firenet-pretrain.h5`
in the terminal.
- (l) Make the scripts executable by typing `chmod +x *.py` in the terminal.
- (3) Start the Docker instance by typing `docker run -u $(id -u):$(id -g) --rm -it -v "$PWD:/tmp" -w /tmp tensorflow:2.9.3` in the terminal.'
- (4) Construct a randomly initialized FireNet model by typing `./firenet.py -a 10 -f selu -i 64 -o firenet-i0.h5 -r 123456789` in the terminal. Answer question (2).
- (5) View the untrained model in netron using the graphical interface or a web browser (<https://netron.app>). Answer question (3).
- (6) Create a directory to save trained models to by typing `mkdir firstTrain` in the terminal.
- (7) Train the model by typing `time ./trainImages.py -a 10 -b 64 -c -e 64 -f ce+clr+aw -i 64 -l 0.01 -m firenet-i0.h5 -o firstTrain -t NA-Acer-64-train.tfr -v NA-Acer-64-validation.tfr` in the terminal. Answer question (4).
- (8) Evaluate the best model from the training by typing `./testImages.py -a 10 -c -i 64 -m $(ls -ltr firstTrain/*/best-model.h5 | awk '{print $9}' | tail -1) -t NA-Acer-64-test.tfr` in the terminal. Answer question (5).
- (9) Create a directory to save trained models to by typing `mkdir secondTrain` in the terminal.
- (10) Fine-tune the pretrained FireNet the model typing `time ./finetuneImages.py -a 10 -b 64 -c -e 64 -f ce+clr+aw -i 64 -l 0.01 -m firenet-pretrain.h5 -o secondTrain -t NA-Acer-64-train.tfr -v NA-Acer-64-validation.tfr` in the terminal. Answer question (6).
- (11) Evaluate the best model from the fine-tuning by typing `./testImages.py -a 10 -c -i 64 -m $(ls -ltr secondTrain/*/best-model.h5 | awk '{print $9}' | tail -1) -t NA-Acer-64-test.tfr` in the terminal. Answer question (7).
- (12) Close the Docker image by typing `exit` in the terminal.

Questions (<https://forms.gle/5A1jaYFK6TKPfJQF8>)

- (1) For task (2)(e), explain what each part of the command does.
- (2) For task (4), explain what each argument of the command does.
- (3) For task (5):
 - (a) How many parameters does this FireNet model have?
 - (b) How many parameters are trainable?
 - (c) What is the size of the output layer?
 - (d) How large of an image will be processed by the model?
- (4) For task (7):

- (a) Explain what each part of the command does.
 - (b) How long (real time == wall clock time) did it take to train the model?
 - (c) What was the model's accuracy, AUPRC, and F_1 on the validation dataset?
 - (d) Did the model overfit during training?
- (5) For task (8):
- (a) What was the model's accuracy, AUPRC, and F_1 on the test dataset?
 - (b) Is this very different from the corresponding values on the validation dataset?
 - (c) Would the model benefit from additional epochs of training?
- (6) For task (10):
- (a) How many parameters does this FireNet model have?
 - (b) How many parameters are trainable?
 - (c) Compare and contrast the number of trainable parameters and their location in this model and in 'firenet-i0.h5'.
 - (d) How long (real time == wall clock time) did it take to fine-tune the model?
 - (e) How does this compare to the first training time?
 - (f) What was the model's accuracy, AUPRC, and F_1 on the validation dataset?
 - (g) Did the model overfit during training?
- (7) For task (11):
- (a) What was the model's accuracy, AUPRC, and F_1 on the test dataset?
 - (b) Which training worked better? Explain.

Literature cited

Gupta, K. D., , D. K. Sharma, S. Ahmed, H. Gupta, D. Gupta & C.-H. Hsu. 2023. A novel lightweight deep learning-based histopathological image classification model for IoMT. *Neural Processing Letters* 55: 205–228.

Iandola, F., S. Han, M. Moskewicz, K. Ashraf, W. Dally & K. Keutzer. 2016. SqueezeNet: AlexNet-level accuracy with 50x fewer parameters and <0.5mb model size. *arXiv 1602.07360*.

Due at the start of class May 2.