

Laboratory 14: TensorFlow DNA sequence regression

In this laboratory exercise, you will use TensorFlow to train an efficient machine learning model to predict protein stability. The dataset is a mixture of naturally occurring miniproteins (43, 46, or 50 amino acid residues) and artificially induced point mutations that have been characterized using a high-throughput assay (Rocklin et al. 2017). The dataset has been partitioned into 11,905 (85.33%) train, 1024 (7.33%) test, and 1024 (7.33%) validate samples.

You will convert the DNA sequences to integer tokens with a 10-bit unigram language model (SentencePiece; Kudo 2018), convert each token into a vector of floating-point numbers with two small embeddings using the quotient-remainder trick (Shi et al. 2020), and then infer protein stability with a small-scale Compact Convolutional Transformer (CCT; Hassani et al. 2021) modified to use the more computationally efficient AveragePooling attention (Yu et al. 2021) in place of conventional attention.

Tasks

- (1) Create a TensorFlow environment to build and train the CCT model.
 - (a) Make a working project directory by typing `mkdir stability && cd stability` in the terminal.
 - (b) Start a Docker build file by typing `echo 'FROM tensorflow/tensorflow:2.12.0' > Dockerfile` in the terminal.
 - (c) Add a RUN statement to the build file by typing `echo 'RUN python3 -m pip install --upgrade pip && python3 -m pip install --upgrade setuptools && python3 -m pip install tensorflow-addons && python3 -m pip install tensorflow-text' >> Dockerfile` in the terminal.
 - (d) Build the Docker image by typing `docker build -t 'tensorflow:2.12.0' .` in the terminal. This process may take several minutes to complete depending on your network speed.
 - (e) Download the segmented data by typing `PARTITIONS=('test' 'train' 'validation'); for PARTITION in "${PARTITIONS[@]}"; do wget 'https://github.com/dpl10/phytoinformatics2023/raw/main/stability-${PARTITION}.tfr'; done` in the terminal. Answer question (1).
 - (f) Download the model script by typing `wget https://raw.githubusercontent.com/dpl10/phytoinformatics2023/main/cct.py` in the terminal.
 - (g) Download the training and testing scripts by typing `TYPES=('finetune' 'predict' 'test' 'train'); for TYPE in "${TYPES[@]}"; do wget 'https://github.com/dpl10/phytoinformatics2023/raw/main/${TYPE}Sequences.py'; done` in the terminal.
 - (h) Make the scripts executable by typing `chmod +x *.py` in the terminal.
 - (i) Download the SentencePeice 10-bit unigram model by typing `wget https://raw.githubusercontent.com/dpl10/phytoinformatics2023/main/unigram-10.pb` in the terminal.
 - (j) Download a pretrained CCT model by typing `wget https://github.com/dpl10/phytoinformatics2023/raw/main/cct-pretrain.h5` in the terminal.
- (2) Start the Docker instance by typing `docker run -u $(id -u):$(id -g) --rm -it -v "$PWD:/tmp" -w /tmp tensorflow:2.12.0` in the terminal.'

- (3) Construct a randomly initialized CCT model by typing `./cct.py -a 1 -b -d 32 -f selu -l 2 -m 4 -o cct.h5 -p 0,1,2,3 -r 123456789 -u 48 -v 10` in the terminal. Answer question (2).
- (4) View the untrained model in netron using the graphical interface or a web browser (<https://netron.app>). Answer question (3).
- (5) Make a directory to save the training result by typing `mkdir denovo` in the terminal.
- (6) Train the CCT model by typing `time ./trainSequences.py -a 1 -b 64 -c -e 64 -f mse+clr+a -l 0.05 -m cct.h5 -o denovo -r 123456789 -s unigram-10.pb -t stability-train.tfr -u 48 -v stability-validation.tfr` in the terminal. Depending upon your computer, this step will take between 2 minutes and 2 hours. Answer question (4).
- (7) Evaluate the best-trained model by typing `./testSequences.py -a 1 -b 64 -c -m $(ls -ltr denovo/*/best-model.h5 | awk '{print $9}' | tail -1) -s unigram-10.pb -t stability-test.tfr -u 48` in the terminal. Answer question (5).
- (8) Create a directory to save finetuned models to by typing `mkdir finetune` in the terminal.
- (9) Fine-tune the pretrained CCT model by typing `time ./finetuneSequences.py -a 1 -b 64 -c -e 64 -f mse+clr+a -l 0.001 -m cct-pretrain.h5 -o finetune -r 123456789 -s unigram-10.pb -t stability-train.tfr -u 48 -v stability-validation.tfr` in the terminal. Answer question (6).
- (10) Evaluate the best model from the fine-tuning by typing `./testSequences.py -a 1 -b 64 -c -m $(ls -ltr finetune/*/best-model.h5 | awk '{print $9}' | tail -1) -s unigram-10.pb -t stability-test.tfr -u 48` in the terminal. Answer question (7).
- (11) Output a table of [-1,+1] rescaled known values and predicted values by typing `./predictSequences.py -a 1 -c -m $(ls -ltr denovo/*/best-model.h5 | awk '{print $9}' | tail -1) -s unigram-10.pb -t stability-test.tfr -u 48 > denovo-predictions.tsv` in the terminal.
- (12) Output a similar table for the finetuned model by typing `./predictSequences.py -a 1 -c -m $(ls -ltr finetune/*/best-model.h5 | awk '{print $9}' | tail -1) -s unigram-10.pb -t stability-test.tfr -u 48 > finetune-predictions.tsv` in the terminal.
- (13) Close the Docker image by typing `exit` in the terminal.
- (14) To evaluate the *de novo* prediction calculations, type `tail -n +2 denovo-predictions.tsv | datamash ppearson 1:2` in the terminal.
- (15) To evaluate the finetune prediction calculations, type `tail -n +2 finetune-predictions.tsv | datamash ppearson 1:2` in the terminal. Answer question (8).

Questions (<https://forms.gle/QwQ4syDwEFZFxQUR6>)

- (1) For task (1)(e), explain what each part of the command does.
- (2) For task (3), explain what each argument of the command does.
- (3) For task (4):

- (a) How many parameters does this CCT model have?
 - (b) How many parameters are trainable?
 - (c) What is the size of each transformer unit?
 - (d) How many sequence tokens per sample can be processed by the model?
 - (e) What is the size of the output layer?
- (4) For task (6):
- (a) Explain what each part of the command does.
 - (b) How long (real time == wall clock time) did it take to train the model?
 - (c) What was the model's MSE on the validation dataset?
 - (d) Did the model overfit during training?
- (5) For task (7):
- (a) What was the model's MSE on the test dataset?
 - (b) Is this very different from the corresponding values on the validation dataset?
 - (c) Would the model benefit from additional epochs of training?
- (6) For task (9):
- (a) How many parameters does this CCT model have?
 - (b) How many parameters are trainable?
 - (c) Compare and contrast the number of trainable parameters and their location in this model and in 'cct.h5'.
 - (d) How long (real time == wall clock time) did it take to fine-tune the model?
 - (e) How does this compare to the *de novo* training time?
 - (f) What was the model's MSE on the validation dataset?
 - (g) Did the model overfit during training?
- (7) For task (10):
- (a) What was the model's MSE on the test dataset?
 - (b) Which training worked better? Explain.
- (8) For task (15):
- (a) Which model is better correlated with the test dataset?
 - (b) Is there any pattern to the errors made by the models?
 - (c) How could you (perhaps) overcome this bias?

Literature cited

- Hassani, A., S. Walton, N. Shah, A. Abuduweili, J. Li & H. Shi.** 2021. Escaping the big data paradigm with compact transformers. arXiv 2104.05704.
- Kudo, T.** 2018. Subword regularization: improving neural network translation models with multiple subword candidates. arXiv 1804.10959.
- Rocklin, G., T. Chidyausiku, I. Goreshnik, A. Ford, S. Houliston, A. Lemak, L. Carter, R. Ravichandran, V. K. Mulligan, A. Chevalier, C. H. Arrowsmith & D. Baker.** 2017. Global analysis of protein folding using massively parallel design, synthesis, and testing. *Science* 357: 168–175.
- Shi, H.-J., D. Mudigere, M. Naumov & J. Yang.** 2020. Compositional embeddings using complementary partitions for memory-efficient recommendation systems. *In: Proceedings of the 26th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*, pp. 165–175. Association for Computing Machinery.
- Yu, M., M. Luo, P. Zhou, C. Si, Y. Zhou, X. Wang, J. Feng & S. Yan.** 2021. MetaFormer is actually what you need for vision. arXiv 2111.11418.

Due at the start of class May 9.